

Nick (Seungheon) Oh

nick.sh.oh@gmail.com

EDUCATION

Apr 2022 — Sept 2023	MSc Data Science and Artificial Intelligence, University of London <i>Distinction</i>	London
Sept 2016 — Jun 2020	BSc Government and Economics, London School of Economics & Political Science (LSE) <i>Upper Second Class Honours</i> [Cumulative GPA: 3.5-7/4.0]	London

WORK EXPERIENCE

Jan 2023 — Present	Researcher, socius labs Founded socius labs , an independent research lab supported by the London School of Economics (LSE) and LSE Generate, bridging human sciences and artificial intelligence through two parallel research streams: applying insights from human sciences to enhance AI systems, and leveraging AI as a platform to run human theories at scale.	London
Mar 2025 — Sept 2026	Defence AI researcher, Ministry of National Defence Assigned to the S/C3 (Command, Control & Communications) Capability Requirements Branch, Training and Doctrine Command (TRADOC); on assignment from the Military Science and Technology Research Unit, Republic of Korea Army Headquarters [received TRADOC Commander's commendation for outstanding performance]	Republic of Korea
Nov 2021 — Oct 2022	Junior NLP engineer, Numen Capital Contributed to Knowsis research team extracting trading signals from non-traditional data sources (social media) through NLP and machine learning. Built transformer-based ESG (Environment, Social, Governance) tweet classifier achieving 97% accuracy, implemented rule-based sentiment analyser with custom financial dictionaries, and integrated these NLP-derived features into deep learning models for market movement forecasting. Developed end-to-end pipeline from raw text to actionable trading signals.	London

SELECTED PUBLICATIONS

Oh, N., & Gobet, F. (2026). Small Foundation Models of Human Cognition and Behaviour. *The 3rd Conference on Language Modeling (COLM'26)* [\[Paper\]](#)

Oh, N., & Jin, H. (2026). Reliability, Faithfulness, and the Limits of Post-hoc Explanations of Opaque Scientific Models. *Philosophy of Machine Learning Workshop @ the 43rd International Conference on Machine Learning (PhilML Workshop, ICML'26)* [\[Paper\]](#)

Oh, N., & Gobet, F. (2025). Monitor-Generate-Verify (MGV): Formalising Metacognitive Theory for Language Model Reasoning. *Workshop on Foundations of Reasoning in Language Models @ the 39th Annual Conference on Neural Information Processing Systems (FoRLM Workshop, NeurIPS'25)* [\[Paper\]](#)

CONFERENCE PROCEEDINGS	Oh, N., & Gobet, F. (2026). Small Foundation Models of Human Cognition and Behaviour. <i>The 3rd Conference on Language Modeling (COLM'26)</i> [Paper] Oh, N., Vrakas, G. D., Brooke, S. J., Morinière, S., & Duke, T. (2025). PETLP: A Privacy-by-Design Pipeline for Social Media Data in AI Research. <i>Proceedings of the 8th AAAI/ACM Conference on AI, Ethics, and Society (AIES'25)</i> [Paper] [Poster] Oh, N. (2024). Sentibank: A Unified Resource of Sentiment Lexicons and Dictionaries. <i>Proceedings of the 18th International AAAI Conference on Web and Social Media (ICWSM'24)</i> [Paper] [Poster]	
WORKSHOPS	Oh, N., & Jin, H. (2026). Reliability, Faithfulness, and the Limits of Post-hoc Explanations of Opaque Scientific Models. <i>Philosophy of Machine Learning Workshop @ the 43rd International Conference on Machine Learning (PhilML Workshop, ICML'26)</i> [Paper] [Poster] Oh, N., & Gobet, F. (2025). Monitor-Generate-Verify (MGV): Formalising Metacognitive Theory for Language Model Reasoning. <i>Workshop on Foundations of Reasoning in Language Models @ the 39th Annual Conference on Neural Information Processing Systems (FoRLM Workshop, NeurIPS'25)</i> [Paper] [Poster] Oh, N. (2025). Before you <think>, monitor: Implementing Flavell's metacognitive framework in LLMs. <i>Workshop on the Application of LLM Explainability to Reasoning and Planning @ Conference on Language Modeling (XLLM-Reason-Plan Workshop, COLM'25)</i> [Paper] [Poster] Oh, N., & Gobet, F. (2024). System 1.5: Designing Metacognition in Artificial Intelligence. <i>System 2 Reasoning at Scale Workshop @ the 38th Annual Conference on Neural Information Processing Systems (Sys2-Reasoning Workshop, NeurIPS'24; Spotlight)</i> [Paper] [Poster] [Spotlight Presentation (Video)] Oh, N. (2024). In Defence of Post-hoc Explainability. <i>Interpretable AI: Past, Present and Future Workshop @ the 38th Annual Conference on Neural Information Processing Systems (IAI Workshop, NeurIPS'24)</i> [Paper] [Poster]	
WORKING PAPERS	High-Intensity Ambivalence: Beyond Acceptance and Resistance to Generative AI in Occupational Communities (in review), with Humphrey, J. & Ross, D. Stochastic Optimal Control for Latent Reasoning in Large Language Models, with Cho, Y., Lee, M., Joo, S. & Lee, A. Application of Steering Vectors in Political Science, with Cho, Y. & Gobet, F.	
TEACHING & TUTORIALS	<i>Scraping Reddit the Right Way: A Guide to Legal and Ethical Data Collection with RedditHarbor</i> — Tutorial @ ICWSM'24, co-organised with the Open Data Institute (ODI)	
EXHIBITIONS	Oh, N., & Park, A. (2026). artefact(s): LeNet-1. <i>Art Gallery @ the IEEE/CVF Conference on Computer Vision and Pattern Recognition, Denver (CVPR'26; IEEE TCPAMI Art Award; Spotlight)</i> [Gallery] [Video]	
RELEVANT TRAINING & SKILLS	LSE MA314: Algorithms and Programming (First Class) LSE GV324: Applied Quantitative Methods for Political Science (First Class) LSE GV249: Research Design in Political Science (Upper Second Class)	
	LSE Methodology & Quantitative Training: Causal inference and natural experiments in the social sciences (regression, matching, panel data, diff-in-diff, instrumental variables, regression discontinuity) Python: PyTorch, spaCy, NLTK, transformers, trl, unsloth, vllm, peft, PRAW, OpenAPI	